

ASSIGNMENT 4 - CLASSIFICATION

The dataset "Breast Cancer" contains various features related to cell measurements, including characteristics such as radius, texture, smoothness, and compactness, along with a target variable indicating whether the tumor is malignant or benign. The primary goal of this project is to design and implement a comprehensive classification system that addresses key challenges in classifying tumors, such as feature scaling, model selection, and handling class imbalance. By applying effective classification algorithms, the objective is to analyze and predict whether a tumor is malignant or benign based on the provided features, ultimately enhancing the overall quality, reliability, and usability of the model for further analysis and machine learning applications. This task will focus on implementing and comparing multiple classification techniques to determine the best model for tumor classification.

SOURCE

The Breast Cancer dataset used for this project is available in the sklearn library. It can be loaded using the `load_breast_cancer()` function from `sklearn.datasets`.

IMPORTING MODULES

```
In [4]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
import sys
if not sys.warnoptions:
    warnings.simplefilter("ignore")
```

LOADING & PREPROCESSING

1. LOAD THE DATA AND CONVERT INTO DATA FRAME

```
In [6]: # LOAD THE DATASET
from sklearn.datasets import load_breast_cancer

data = load_breast_cancer()

# Convert to DataFrame
df = pd.DataFrame(data.data, columns=data.feature_names)
df['target'] = data.target

# Display the first few rows
print(df.head())
```

| | mean radius | mean texture | mean perimeter | mean area | mean smoothness | \ |
|---|------------------------|-------------------------|---------------------|----------------------|-----------------|---|
| 0 | 17.99 | 10.38 | 122.80 | 1001.0 | 0.11840 | |
| 1 | 20.57 | 17.77 | 132.90 | 1326.0 | 0.08474 | |
| 2 | 19.69 | 21.25 | 130.00 | 1203.0 | 0.10960 | |
| 3 | 11.42 | 20.38 | 77.58 | 386.1 | 0.14250 | |
| 4 | 20.29 | 14.34 | 135.10 | 1297.0 | 0.10030 | |
| | mean compactness | mean concavity | mean concave points | mean symmetry | \ | |
| 0 | 0.27760 | 0.3001 | 0.14710 | 0.2419 | | |
| 1 | 0.07864 | 0.0869 | 0.07017 | 0.1812 | | |
| 2 | 0.15990 | 0.1974 | 0.12790 | 0.2069 | | |
| 3 | 0.28390 | 0.2414 | 0.10520 | 0.2597 | | |
| 4 | 0.13280 | 0.1980 | 0.10430 | 0.1809 | | |
| | mean fractal dimension | ... | worst texture | worst perimeter | worst area | \ |
| 0 | 0.07871 | ... | 17.33 | 184.60 | 2019.0 | |
| 1 | 0.05667 | ... | 23.41 | 158.80 | 1956.0 | |
| 2 | 0.05999 | ... | 25.53 | 152.50 | 1709.0 | |
| 3 | 0.09744 | ... | 26.50 | 98.87 | 567.7 | |
| 4 | 0.05883 | ... | 16.67 | 152.20 | 1575.0 | |
| | worst smoothness | worst compactness | worst concavity | worst concave points | \ | |
| 0 | 0.1622 | 0.6656 | 0.7119 | 0.2654 | | |
| 1 | 0.1238 | 0.1866 | 0.2416 | 0.1860 | | |
| 2 | 0.1444 | 0.4245 | 0.4504 | 0.2430 | | |
| 3 | 0.2098 | 0.8663 | 0.6869 | 0.2575 | | |
| 4 | 0.1374 | 0.2050 | 0.4000 | 0.1625 | | |
| | worst symmetry | worst fractal dimension | target | | | |
| 0 | 0.4601 | 0.11890 | 0 | | | |
| 1 | 0.2750 | 0.08902 | 0 | | | |
| 2 | 0.3613 | 0.08758 | 0 | | | |
| 3 | 0.6638 | 0.17300 | 0 | | | |
| 4 | 0.2364 | 0.07678 | 0 | | | |

[5 rows x 31 columns]

2. DISPLAY FIRST & LAST ROWS

```
In [12]: # DISPLAY FIRST FEW ROWS TO UNDERSTAND THE STRUCTURE OF THE DATA
print(df.head())
```

| | mean radius | mean texture | mean perimeter | mean area | mean smoothness | \ |
|---|-------------|--------------|----------------|-----------|-----------------|---|
| 0 | 17.99 | 10.38 | 122.80 | 1001.0 | 0.11840 | |
| 1 | 20.57 | 17.77 | 132.90 | 1326.0 | 0.08474 | |
| 2 | 19.69 | 21.25 | 130.00 | 1203.0 | 0.10960 | |
| 3 | 11.42 | 20.38 | 77.58 | 386.1 | 0.14250 | |
| 4 | 20.29 | 14.34 | 135.10 | 1297.0 | 0.10030 | |

| | mean compactness | mean concavity | mean concave points | mean symmetry | \ |
|---|------------------|----------------|---------------------|---------------|---|
| 0 | 0.27760 | 0.3001 | 0.14710 | 0.2419 | |
| 1 | 0.07864 | 0.0869 | 0.07017 | 0.1812 | |
| 2 | 0.15990 | 0.1974 | 0.12790 | 0.2069 | |
| 3 | 0.28390 | 0.2414 | 0.10520 | 0.2597 | |
| 4 | 0.13280 | 0.1980 | 0.10430 | 0.1809 | |

| | mean fractal dimension | ... | worst texture | worst perimeter | worst area | \ |
|---|------------------------|-----|---------------|-----------------|------------|---|
| 0 | 0.07871 | ... | 17.33 | 184.60 | 2019.0 | |
| 1 | 0.05667 | ... | 23.41 | 158.80 | 1956.0 | |
| 2 | 0.05999 | ... | 25.53 | 152.50 | 1709.0 | |
| 3 | 0.09744 | ... | 26.50 | 98.87 | 567.7 | |
| 4 | 0.05883 | ... | 16.67 | 152.20 | 1575.0 | |

| | worst smoothness | worst compactness | worst concavity | worst concave points | \ |
|---|------------------|-------------------|-----------------|----------------------|---|
| 0 | 0.1622 | 0.6656 | 0.7119 | 0.2654 | |
| 1 | 0.1238 | 0.1866 | 0.2416 | 0.1860 | |
| 2 | 0.1444 | 0.4245 | 0.4504 | 0.2430 | |
| 3 | 0.2098 | 0.8663 | 0.6869 | 0.2575 | |
| 4 | 0.1374 | 0.2050 | 0.4000 | 0.1625 | |

| | worst symmetry | worst fractal dimension | target |
|---|----------------|-------------------------|--------|
| 0 | 0.4601 | 0.11890 | 0 |
| 1 | 0.2750 | 0.08902 | 0 |
| 2 | 0.3613 | 0.08758 | 0 |
| 3 | 0.6638 | 0.17300 | 0 |
| 4 | 0.2364 | 0.07678 | 0 |

[5 rows x 31 columns]

```
In [10]: # DISPLAY LAST FEW ROWS TO UNDERSTAND THE STRUCTURE OF THE DATA
```

```
print(df.tail())
```

| | mean radius | mean texture | mean perimeter | mean area | mean smoothness | \ |
|-----|-------------|--------------|----------------|-----------|-----------------|---|
| 564 | 21.56 | 22.39 | 142.00 | 1479.0 | 0.11100 | |
| 565 | 20.13 | 28.25 | 131.20 | 1261.0 | 0.09780 | |
| 566 | 16.60 | 28.08 | 108.30 | 858.1 | 0.08455 | |
| 567 | 20.60 | 29.33 | 140.10 | 1265.0 | 0.11780 | |
| 568 | 7.76 | 24.54 | 47.92 | 181.0 | 0.05263 | |

| | mean compactness | mean concavity | mean concave points | mean symmetry | \ |
|-----|------------------|----------------|---------------------|---------------|---|
| 564 | 0.11590 | 0.24390 | 0.13890 | 0.1726 | |
| 565 | 0.10340 | 0.14400 | 0.09791 | 0.1752 | |
| 566 | 0.10230 | 0.09251 | 0.05302 | 0.1590 | |
| 567 | 0.27700 | 0.35140 | 0.15200 | 0.2397 | |
| 568 | 0.04362 | 0.00000 | 0.00000 | 0.1587 | |

| | mean fractal dimension | ... | worst texture | worst perimeter | worst area | \ |
|-----|------------------------|-----|---------------|-----------------|------------|---|
| 564 | 0.05623 | ... | 26.40 | 166.10 | 2027.0 | |
| 565 | 0.05533 | ... | 38.25 | 155.00 | 1731.0 | |
| 566 | 0.05648 | ... | 34.12 | 126.70 | 1124.0 | |
| 567 | 0.07016 | ... | 39.42 | 184.60 | 1821.0 | |
| 568 | 0.05884 | ... | 30.37 | 59.16 | 268.6 | |

| | worst smoothness | worst compactness | worst concavity | \ |
|-----|------------------|-------------------|-----------------|---|
| 564 | 0.14100 | 0.21130 | 0.4107 | |
| 565 | 0.11660 | 0.19220 | 0.3215 | |
| 566 | 0.11390 | 0.30940 | 0.3403 | |
| 567 | 0.16500 | 0.86810 | 0.9387 | |
| 568 | 0.08996 | 0.06444 | 0.0000 | |

| | worst concave points | worst symmetry | worst fractal dimension | target |
|-----|----------------------|----------------|-------------------------|--------|
| 564 | 0.2216 | 0.2060 | 0.07115 | 0 |
| 565 | 0.1628 | 0.2572 | 0.06637 | 0 |
| 566 | 0.1418 | 0.2218 | 0.07820 | 0 |
| 567 | 0.2650 | 0.4087 | 0.12400 | 0 |
| 568 | 0.0000 | 0.2871 | 0.07039 | 1 |

```
[5 rows x 31 columns]
```

3. DATATYPE OF EACH COLUMN

```
In [14]: # DISPLAY DATA TYPE OF EACH COLUMN  
print("Dataset Info:")  
df.info()
```

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 569 entries, 0 to 568

Data columns (total 31 columns):

| # | Column | Non-Null Count | Dtype |
|----|-------------------------|----------------|---------|
| 0 | mean radius | 569 non-null | float64 |
| 1 | mean texture | 569 non-null | float64 |
| 2 | mean perimeter | 569 non-null | float64 |
| 3 | mean area | 569 non-null | float64 |
| 4 | mean smoothness | 569 non-null | float64 |
| 5 | mean compactness | 569 non-null | float64 |
| 6 | mean concavity | 569 non-null | float64 |
| 7 | mean concave points | 569 non-null | float64 |
| 8 | mean symmetry | 569 non-null | float64 |
| 9 | mean fractal dimension | 569 non-null | float64 |
| 10 | radius error | 569 non-null | float64 |
| 11 | texture error | 569 non-null | float64 |
| 12 | perimeter error | 569 non-null | float64 |
| 13 | area error | 569 non-null | float64 |
| 14 | smoothness error | 569 non-null | float64 |
| 15 | compactness error | 569 non-null | float64 |
| 16 | concavity error | 569 non-null | float64 |
| 17 | concave points error | 569 non-null | float64 |
| 18 | symmetry error | 569 non-null | float64 |
| 19 | fractal dimension error | 569 non-null | float64 |
| 20 | worst radius | 569 non-null | float64 |
| 21 | worst texture | 569 non-null | float64 |
| 22 | worst perimeter | 569 non-null | float64 |
| 23 | worst area | 569 non-null | float64 |
| 24 | worst smoothness | 569 non-null | float64 |
| 25 | worst compactness | 569 non-null | float64 |
| 26 | worst concavity | 569 non-null | float64 |
| 27 | worst concave points | 569 non-null | float64 |
| 28 | worst symmetry | 569 non-null | float64 |
| 29 | worst fractal dimension | 569 non-null | float64 |
| 30 | target | 569 non-null | int32 |

dtypes: float64(30), int32(1)

memory usage: 135.7 KB

4. STATISTICAL SUMMARY OF DATA

```
In [16]: # DISPLAY STATISTICAL SUMMARY
print("Statistical Summary:")
df.describe()
```

Statistical Summary:

Out[16]:

| | mean radius | mean texture | mean perimeter | mean area | mean smoothness | mean compactness | mean concavity | mean concave points | mean symmetry | me frac dimens |
|--------------|----------------|-----------------|-------------------|-------------|--------------------|---------------------|-------------------|---------------------------|------------------|----------------------|
| count | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 569.000000 |
| mean | 14.127292 | 19.289649 | 91.969033 | 654.889104 | 0.096360 | 0.104341 | 0.088799 | 0.048919 | 0.181162 | 0.0627 |
| std | 3.524049 | 4.301036 | 24.298981 | 351.914129 | 0.014064 | 0.052813 | 0.079720 | 0.038803 | 0.027414 | 0.0070 |
| min | 6.981000 | 9.710000 | 43.790000 | 143.500000 | 0.052630 | 0.019380 | 0.000000 | 0.000000 | 0.106000 | 0.0495 |
| 25% | 11.700000 | 16.170000 | 75.170000 | 420.300000 | 0.086370 | 0.064920 | 0.029560 | 0.020310 | 0.161900 | 0.0577 |
| 50% | 13.370000 | 18.840000 | 86.240000 | 551.100000 | 0.095870 | 0.092630 | 0.061540 | 0.033500 | 0.179200 | 0.0611 |
| 75% | 15.780000 | 21.800000 | 104.100000 | 782.700000 | 0.105300 | 0.130400 | 0.130700 | 0.074000 | 0.195700 | 0.0667 |
| max | 28.110000 | 39.280000 | 188.500000 | 2501.000000 | 0.163400 | 0.345400 | 0.426800 | 0.201200 | 0.304000 | 0.0974 |

8 rows × 11 columns



5. DISPLAY ALL COLUMN NAMES

```
In [18]: # DISPLAY PARTICULAR COLUMN
print("Columns of the dataset:")
df.columns
```

Columns of the dataset:


```
Out[18]: Index(['mean radius', 'mean texture', 'mean perimeter', 'mean area',  
              'mean smoothness', 'mean compactness', 'mean concavity',  
              'mean concave points', 'mean symmetry', 'mean fractal dimension',  
              'radius error', 'texture error', 'perimeter error', 'area error',  
              'smoothness error', 'compactness error', 'concavity error',  
              'concave points error', 'symmetry error', 'fractal dimension error',  
              'worst radius', 'worst texture', 'worst perimeter', 'worst area',  
              'worst smoothness', 'worst compactness', 'worst concavity',  
              'worst concave points', 'worst symmetry', 'worst fractal dimension',  
              'target'],  
             dtype='object')
```

6. NULL / MISSING VALUES IN EACH COLUMN

```
In [20]: # DISPLAY NULL VALUES IN EACH COLUMN  
print("Null values in each column:")  
print(df.isnull().sum())
```

```
Null values in each column:
mean radius      0
mean texture     0
mean perimeter   0
mean area        0
mean smoothness  0
mean compactness 0
mean concavity   0
mean concave points 0
mean symmetry    0
mean fractal dimension 0
radius error     0
texture error    0
perimeter error  0
area error       0
smoothness error 0
compactness error 0
concavity error  0
concave points error 0
symmetry error   0
fractal dimension error 0
worst radius     0
worst texture    0
worst perimeter  0
worst area       0
worst smoothness 0
worst compactness 0
worst concavity  0
worst concave points 0
worst symmetry   0
worst fractal dimension 0
target          0
dtype: int64
```

7. DUPLICATE VALUES

```
In [22]: # FINDING THE TOTAL NO OF DUPLICATES
df.duplicated().sum()
```

```
Out[22]: 0
```

8. FEATURE SCALING

```
In [26]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

X = df.drop(columns=['target'])
y = df['target']

X_scaled = scaler.fit_transform(X)

print("\nScaled Feature Data (First 5 rows):")
print(X_scaled[:5])
```

Scaled Feature Data (First 5 rows):

```
[[ 1.09706398e+00 -2.07333501e+00 1.26993369e+00 9.84374905e-01
  1.56846633e+00 3.28351467e+00 2.65287398e+00 2.53247522e+00
  2.21751501e+00 2.25574689e+00 2.48973393e+00 -5.65265059e-01
  2.83303087e+00 2.48757756e+00 -2.14001647e-01 1.31686157e+00
  7.24026158e-01 6.60819941e-01 1.14875667e+00 9.07083081e-01
  1.88668963e+00 -1.35929347e+00 2.30360062e+00 2.00123749e+00
  1.30768627e+00 2.61666502e+00 2.10952635e+00 2.29607613e+00
  2.75062224e+00 1.93701461e+00]
 [ 1.82982061e+00 -3.53632408e-01 1.68595471e+00 1.90870825e+00
 -8.26962447e-01 -4.87071673e-01 -2.38458552e-02 5.48144156e-01
  1.39236330e-03 -8.68652457e-01 4.99254601e-01 -8.76243603e-01
  2.63326966e-01 7.42401948e-01 -6.05350847e-01 -6.92926270e-01
 -4.40780058e-01 2.60162067e-01 -8.05450380e-01 -9.94437403e-02
  1.80592744e+00 -3.69203222e-01 1.53512599e+00 1.89048899e+00
 -3.75611957e-01 -4.30444219e-01 -1.46748968e-01 1.08708430e+00
 -2.43889668e-01 2.81189987e-01]
 [ 1.57988811e+00 4.56186952e-01 1.56650313e+00 1.55888363e+00
 9.42210440e-01 1.05292554e+00 1.36347845e+00 2.03723076e+00
 9.39684817e-01 -3.98007910e-01 1.22867595e+00 -7.80083377e-01
 8.50928301e-01 1.18133606e+00 -2.97005012e-01 8.14973504e-01
 2.13076435e-01 1.42482747e+00 2.37035535e-01 2.93559404e-01
 1.51187025e+00 -2.39743838e-02 1.34747521e+00 1.45628455e+00
 5.27407405e-01 1.08293217e+00 8.54973944e-01 1.95500035e+00
 1.15225500e+00 2.01391209e-01]
 [-7.68909287e-01 2.53732112e-01 -5.92687167e-01 -7.64463792e-01
 3.28355348e+00 3.40290899e+00 1.91589718e+00 1.45170736e+00
 2.86738293e+00 4.91091929e+00 3.26373441e-01 -1.10409044e-01
 2.86593405e-01 -2.88378148e-01 6.89701660e-01 2.74428041e+00
 8.19518384e-01 1.11500701e+00 4.73268037e+00 2.04751088e+00
 -2.81464464e-01 1.33984094e-01 -2.49939304e-01 -5.50021228e-01
 3.39427470e+00 3.89339743e+00 1.98958826e+00 2.17578601e+00
 6.04604135e+00 4.93501034e+00]
 [ 1.75029663e+00 -1.15181643e+00 1.77657315e+00 1.82622928e+00
 2.80371830e-01 5.39340452e-01 1.37101143e+00 1.42849277e+00
 -9.56046689e-03 -5.62449981e-01 1.27054278e+00 -7.90243702e-01
 1.27318941e+00 1.19035676e+00 1.48306716e+00 -4.85198799e-02
 8.28470780e-01 1.14420474e+00 -3.61092272e-01 4.99328134e-01
 1.29857524e+00 -1.46677038e+00 1.33853946e+00 1.22072425e+00
 2.20556166e-01 -3.13394511e-01 6.13178758e-01 7.29259257e-01
 -8.68352984e-01 -3.97099619e-01]]
```

9. SPLITTING THE DATA INTO TRAINING AND TESTING SET

```
In [28]: from sklearn.model_selection import train_test_split

# Split data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.3, random_state=42)

# Display the shape of the split data
print(f"Training data shape: {X_train.shape}")
print(f"Testing data shape: {X_test.shape}")
```

Training data shape: (398, 30)

Testing data shape: (171, 30)

Preprocessing steps with explanations:

1. *Load the Data:*

- The dataset is loaded using `load_breast_cancer()` from `sklearn.datasets`. This gives us the feature data (X) and the target data (y), which indicates whether a tumor is malignant (1) or benign (0).

2. *Convert to DataFrame:*

- Converted the data into pandas DataFrames for easier manipulation and analysis.

3. *Display First and Last Rows:*

- Displayed the first few rows to understand the data structure and confirm it loaded correctly.

4. *Check Data Types:*

- Used `info()` to check the data types of the columns and ensure they are as expected (numerical values).

5. *Statistical Summary:*

- Used `describe()` to view statistics (mean, min, max, etc.) of both features and target to understand their distribution.

6. *Display Column Names:*

- Printed the column names of the features to know what variables we are working with.

7. *Check for Missing Values:*

- Checked for missing values with `isnull().sum()` to ensure the dataset is complete.

8. *Find Duplicate Rows:*

- Checked for duplicate rows using `duplicated().sum()` to ensure there are no repeated records.

9. *Feature Scaling:*

- Scaled the features using `StandardScaler` to ensure that all features are on the same scale, which is important for some machine learning models.

10. *Train-Test Split:*

- Split the data into training and testing sets to evaluate the model's performance on unseen data.

These steps are necessary to clean and prepare the data for better model performance.

CLASSIFICATION ALGORITHMS

IMPLEMENTATION

1. LINEAR REGRESSION ALGORITHM

```
In [37]: from sklearn.linear_model import LogisticRegression
```

```
lr_model = LogisticRegression()
```

```
lr_model.fit(X_train, y_train)
```

```
y_pred = lr_model.predict(X_test)  
y_pred
```

```
Out[37]: array([1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1,  
                0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1,  
                1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1,  
                0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0,  
                1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1,  
                0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0,  
                1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1,  
                1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1])
```

Logistic Regression works by estimating the probability of a binary outcome based on input features, using a logistic function. It assumes a linear relationship between

the features (such as mean radius, texture, smoothness, etc.) and the target variable (malignant or benign tumor). This model is suitable for the Breast Cancer dataset because factors like cell characteristics likely have a linear influence on the likelihood of a tumor being malignant or benign, making it an appropriate choice for classifying tumor types.

2. DECISION TREE CLASSIFIER ALGORITHM

```
In [44]: from sklearn.tree import DecisionTreeClassifier
```

```
dt_model = DecisionTreeClassifier()
```

```
dt_model.fit(X_train, y_train)
```

```
y_pred_dt = dt_model.predict(X_test)
```

```
y_pred_dt
```

```
Out[44]: array([1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1,
        0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1,
        1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1,
        0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0,
        1, 1, 1, 1, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1,
        0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0,
        0, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1,
        0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1])
```

The Decision Tree Classifier works by recursively splitting the data based on feature values to maximize information gain and minimize impurity within each subset. It does not assume a linear relationship between the target variable (malignant or benign tumor) and the input features (such as radius, texture, smoothness, etc.).

This model is suitable for the Breast Cancer dataset because it can capture non-linear relationships and complex interactions between features, such as how various cell characteristics might jointly influence the likelihood of a tumor being malignant or benign in ways that a linear model cannot.

3. RANDOM FOREST CLASSIFIER ALGORITHM

```
In [55]: from sklearn.ensemble import RandomForestClassifier

rf_model = RandomForestClassifier(n_estimators=100, random_state=42)

rf_model.fit(X_train, y_train)

y_pred_rf = rf_model.predict(X_test)
y_pred_rf
```

```
Out[55]: array([1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1,
        0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1,
        1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1,
        0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0,
        1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1,
        0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0,
        1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1,
        1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1])
```

The Random Forest Classifier is an ensemble learning method that constructs multiple decision trees and combines their predictions to enhance accuracy and reduce overfitting. It effectively handles complex, non-linear relationships between features and the target variable. This makes it suitable for the Breast Cancer dataset, as it can capture intricate interactions between factors like cell radius, texture,

smoothness, and compactness, while providing robust predictions and insights into feature importance for classifying tumors as malignant or benign.

4. K NEAREST NEIGHBOUR CLASSIFIER ALGORITHM

```
In [62]: from sklearn.neighbors import KNeighborsClassifier

knn_model = KNeighborsClassifier(n_neighbors=5)

knn_model.fit(X_train, y_train)

y_pred_knn = knn_model.predict(X_test)
y_pred_knn
```

```
Out[62]: array([1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1,
        0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1,
        1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1,
        0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0,
        1, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1,
        0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0,
        1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1,
        1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1])
```

The k-Nearest Neighbors (k-NN) Classifier is a simple, instance-based learning algorithm that classifies data points based on the majority class of their nearest neighbors. It computes the distance between the input data and other points in the feature space to make predictions. This method is well-suited for the Breast Cancer dataset as it can effectively capture complex, non-linear relationships between features like cell radius, texture, and smoothness, which are important for

classifying tumors as malignant or benign. The flexibility of k-NN to handle varied data patterns makes it a robust choice for classification tasks in this dataset.

5. SUPPORT VECTOR CLASSIFIER ALGORITHM

```
In [68]: from sklearn.svm import SVC

svc_model = SVC(kernel='linear')

svc_model.fit(X_train, y_train)

y_pred_svc = svc_model.predict(X_test)
y_pred_svc
```

```
Out[68]: array([1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1,
                0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1,
                1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1,
                0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0,
                1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1,
                0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0,
                1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1,
                1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1])
```

The Support Vector Classifier (SVC) is a classification model that finds a hyperplane that best separates the data into different classes, focusing on maximizing the margin between data points of different classes. It can handle both linear and non-linear decision boundaries by applying kernel functions, such as the Radial Basis Function (RBF). SVC is suitable for the Breast Cancer dataset because it can effectively capture complex, non-linear relationships between features like cell

texture, radius, and smoothness, while also being robust to outliers and effective in high-dimensional spaces.

MODEL EVALUATION

1. LOGISTIC REGRESSION MODEL EVALUATION

```
In [73]: from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# Calculate the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Logistic Regression Accuracy: {accuracy * 100:.2f}%")

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:")
print(cm)

# Classification Report
report = classification_report(y_test, y_pred)
print("\nClassification Report:")
print(report)
```

Logistic Regression Accuracy: 98.25%

Confusion Matrix:

```
[[ 62  1]
 [ 2 106]]
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.97 | 0.98 | 0.98 | 63 |
| 1 | 0.99 | 0.98 | 0.99 | 108 |
| accuracy | | | 0.98 | 171 |
| macro avg | 0.98 | 0.98 | 0.98 | 171 |
| weighted avg | 0.98 | 0.98 | 0.98 | 171 |

2. DECISION TREE CLASSIFIER MODEL EVALUATION

In [47]: `from sklearn.metrics import accuracy_score, confusion_matrix, classification_report`

```
# Calculate the accuracy of the model
accuracy_dtc = accuracy_score(y_test, y_pred_dt)
print(f"Decision Tree Classifier Accuracy: {accuracy_dtc * 100:.2f}%")

# Confusion Matrix
cm_dtc = confusion_matrix(y_test, y_pred_dt)
print("\nConfusion Matrix:")
print(cm_dtc)

# Classification Report
report_dtc = classification_report(y_test, y_pred_dt)
print("\nClassification Report:")
print(report_dtc)
```

Decision Tree Classifier Accuracy: 92.98%

Confusion Matrix:

```
[[60  3]
 [ 9 99]]
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.87 | 0.95 | 0.91 | 63 |
| 1 | 0.97 | 0.92 | 0.94 | 108 |
| accuracy | | | 0.93 | 171 |
| macro avg | 0.92 | 0.93 | 0.93 | 171 |
| weighted avg | 0.93 | 0.93 | 0.93 | 171 |

3. RANDOM FOREST CLASSIFIER MODEL EVALUATION

In [75]: `from sklearn.metrics import accuracy_score, confusion_matrix, classification_report`

```
# Calculate the accuracy of the model
accuracy_rfc = accuracy_score(y_test, y_pred_rf)
print(f"Random Forest Classifier Accuracy: {accuracy_rfc * 100:.2f}%")

# Confusion Matrix
cm_rfc = confusion_matrix(y_test, y_pred_rf)
print("\nConfusion Matrix:")
print(cm_rfc)

# Classification Report
report_rfc = classification_report(y_test, y_pred_rf)
print("\nClassification Report:")
print(report_rfc)
```

Random Forest Classifier Accuracy: 97.08%

Confusion Matrix:

```
[[ 59   4]
 [   1 107]]
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.98 | 0.94 | 0.96 | 63 |
| 1 | 0.96 | 0.99 | 0.98 | 108 |
| accuracy | | | 0.97 | 171 |
| macro avg | 0.97 | 0.96 | 0.97 | 171 |
| weighted avg | 0.97 | 0.97 | 0.97 | 171 |

4. K NEAREST NEIGHBOUR CLASSIFIER MODEL EVALUATION

In [77]: `from sklearn.metrics import accuracy_score, confusion_matrix, classification_report`

```
# Calculate the accuracy of the model
accuracy_knn = accuracy_score(y_test, y_pred_knn)
print(f"KNN Classifier Accuracy: {accuracy_knn * 100:.2f}%")

# Confusion Matrix
cm_knn = confusion_matrix(y_test, y_pred_knn)
print("\nConfusion Matrix:")
print(cm_knn)

# Classification Report
report_knn = classification_report(y_test, y_pred_knn)
print("\nClassification Report:")
print(report_knn)
```

KNN Classifier Accuracy: 95.91%

Confusion Matrix:

```
[[ 59   4]
 [   3 105]]
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.95 | 0.94 | 0.94 | 63 |
| 1 | 0.96 | 0.97 | 0.97 | 108 |
| accuracy | | | 0.96 | 171 |
| macro avg | 0.96 | 0.95 | 0.96 | 171 |
| weighted avg | 0.96 | 0.96 | 0.96 | 171 |

5. SUPPORT VECTOR CLASSIFIER MODEL EVALUATION

In [79]: `from sklearn.metrics import accuracy_score, confusion_matrix, classification_report`

```
# Calculate the accuracy of the model
accuracy_svc = accuracy_score(y_test, y_pred_svc)
print(f"Support Vector Classifier Accuracy: {accuracy_svc * 100:.2f}%")

# Confusion Matrix
cm_svc = confusion_matrix(y_test, y_pred_svc)
print("\nConfusion Matrix:")
print(cm_svc)

# Classification Report
report_svc = classification_report(y_test, y_pred_svc)
print("\nClassification Report:")
print(report_svc)
```


Support Vector Classifier Accuracy: 97.66%

Confusion Matrix:

```
[[ 61   2]
 [   2 106]]
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.97 | 0.97 | 0.97 | 63 |
| 1 | 0.98 | 0.98 | 0.98 | 108 |
| accuracy | | | 0.98 | 171 |
| macro avg | 0.97 | 0.97 | 0.97 | 171 |
| weighted avg | 0.98 | 0.98 | 0.98 | 171 |

Summary of Best and Worst-Performing Models

Best-Performing Model:

The Random Forest Classifier is the best-performing model with the highest accuracy (97.08%) and the best balance in F1-scores across both classes. This indicates that it provides the most accurate predictions and captures complex, non-linear relationships within the data effectively.

Worst-Performing Model:

The Decision Tree Classifier performs the worst among the models tested, with the lowest accuracy (92.98%) and relatively lower F1-scores, particularly for class 0.

While it performs decently for class 1, its overall accuracy and precision/recall balance are weaker compared to the other models.

CONCLUSION

The Random Forest Classifier is the best model for the Breast Cancer dataset because it performs well and effectively handles the complex, non-linear relationships between the features. In contrast, Logistic Regression is the least effective model, likely due to its assumption of a linear decision boundary, which does not capture the intricate patterns within the data.

In []: